

# Chapter 3: Convolution Neural Network



deeplearning.ai

# Convolutional Neural Networks

---

## Computer vision

# Computer Vision Problems

Image Classification



64x64

→ Cat? (0/1)

Object detection


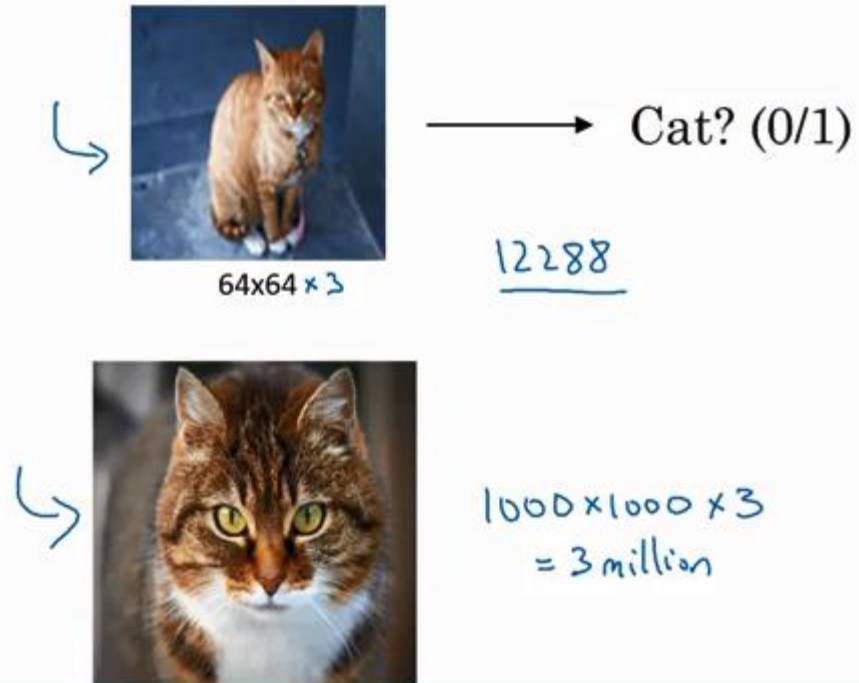


Neural Style Transfer ↙



Andrew Ng

# Deep Learning on large images





deeplearning.ai



# Convolutional Neural Networks

---

## Edge detection example

---

# Computer Vision Problem



Andrew Ng

# Vertical edge detection

$$3 \times 1 + 1 \times 1 + 2 \times 1 + 0 \times 0 + 5 \times 0 + 7 \times 0 + 1 \times -1 + 8 \times -1 + 2 \times -1 = -5$$

3	<u>0</u>	<u>1</u>	<u>2</u>	<u>7</u>	<u>4</u>
1	<u>5</u>	<u>8</u>	<u>9</u>	<u>3</u>	<u>1</u>
2	<u>7</u>	<u>2</u>	<u>5</u>	<u>1</u>	<u>3</u>
0	1	3	1	7	8
4	2	1	6	2	8
2	4	5	2	3	9

6x6

"convolution"  
\*

1	0	-1
1	0	-1
1	0	-1

3x3  
filter

=

<u>-5</u>	<u>-4</u>	0	<u>8</u>
<u>-10</u>	-2	2	3
0	-2	-4	-7
-3	-2	-3	<u>-16</u>


4x4

python: conv-forward  
tensorflow: tf.nn.conv2d  
keras: Conv2D

# Vertical edge detection

$$\begin{array}{|c|c|c|c|c|c|}
 \hline
 10 & 10 & 10 & 0 & 0 & 0 \\
 \hline
 10 & 10 & 10 & 0 & 0 & 0 \\
 \hline
 10 & 10 & 10 & 0 & 0 & 0 \\
 \hline
 10 & 10 & 10 & 0 & 0 & 0 \\
 \hline
 10 & 10 & 10 & 0 & 0 & 0 \\
 \hline
 10 & 10 & 10 & 0 & 0 & 0 \\
 \hline
 \end{array}$$

6x6



\*

$$\begin{array}{|c|c|c|}
 \hline
 1 & 0 & -1 \\
 \hline
 1 & 0 & -1 \\
 \hline
 1 & 0 & -1 \\
 \hline
 \end{array}$$

3x3


\*



=

$$\begin{array}{|c|c|c|c|}
 \hline
 0 & 30 & 30 & 0 \\
 \hline
 0 & 30 & 30 & 0 \\
 \hline
 0 & 30 & 30 & 0 \\
 \hline
 0 & 30 & 30 & 0 \\
 \hline
 \end{array}$$

4x4



Andrew Ng





deeplearning.ai

# Convolutional Neural Networks

---

More edge  
detection

---

## Vertical edge detection examples

10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0



0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10



\*

1	0	-1
1	0	-1
1	0	-1

=

0	30	30	0
0	30	30	0
0	30	30	0
0	30	30	0

\*

1	0	-1
1	0	-1
1	0	-1

=

0	-30	-30	0
0	-30	-30	0
0	-30	-30	0
0	-30	-30	0

Andrew Ng

# Vertical and Horizontal Edge Detection

→

1	0	-1
1	0	-1
1	0	-1

Vertical

→

1	1	1
0	0	0
-1	-1	-1

Horizontal

10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10

6x6

\*

1	1	1
0	0	0
-1	-1	-1

=

0	0	0	0
30	10	-10	-30
30	10	-10	-30
0	0	0	0

# Learning to detect edges

1	0	-1
1	0	-1
1	0	-1



3	0	1	2	7	4
1	5	8	9	3	1
2	7	2	5	1	3
0	1	3	1	7	8
4	2	1	6	2	8
2	4	5	2	3	9

→

1	0	-1
2	0	-2
1	0	-1

Sobel filter



convolution  
\*

$W_1$	$W_2$	$W_3$
$W_4$	$W_5$	$W_6$
$W_7$	$W_8$	$W_9$

3x3

=  
45°  
70°  
73°

3	0	-3
10	0	-10
3	0	-3

Scharr filter




Andrew Ng



deeplearning.ai


# Convolutional Neural Networks

---


## Padding

# Padding

- shrinky output
- throw away info from edge




\*



3x3  
f x f

=



6x6

6x6 → 8x8  
n x n

$n - f + 1 \times n - f + 1$   
 $6 - 3 + 1 = 4$

$p = \text{padding} = \underline{1}$

$n + 2p - f + 1 \times n + 2p - f + 1$   
 $6 + 2 - 3 + 1 \times \underline{\underline{4}} = 6 \times 6$

Andrew Ng

## Valid and Same convolutions

→ no padding

“Valid”:  $n \times n \quad * \quad f \times f \quad \rightarrow \quad \underline{n - f + 1} \times n - f + 1$   
 $6 \times 6 \quad * \quad 3 \times 3 \quad \rightarrow \quad 4 \times 4$

“Same”: Pad so that output size is the same as the input size.

$$n + 2p - f + 1 \times n + 2p - f + 1$$


$$\cancel{n + 2p - f + 1} = \cancel{n} \Rightarrow \boxed{p = \frac{f-1}{2}}$$

$$3 \times 3 \quad p = \frac{3-1}{2} = 1$$

$$5 \times 5 \quad f=5$$

$$p=2$$

$f$  is usually odd



1x1  
3x3  
5x5  
7x7



deeplearning.ai

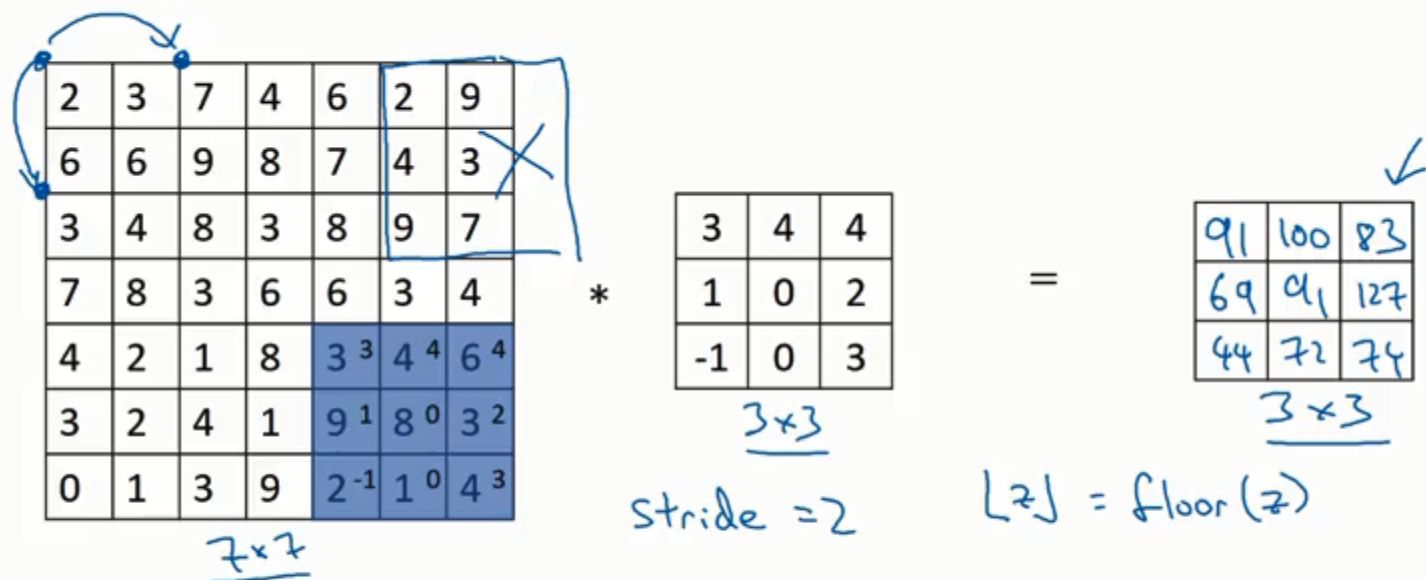
# Convolutional Neural Networks

---

## Strided convolutions



# Strided convolution



$$\begin{array}{l}
 n \times n \quad * \quad f \times f \\
 \text{padding } p \quad \text{stride } s \\
 s = 2
 \end{array}$$

$$\left\lfloor \frac{n + 2p - f}{s} + 1 \right\rfloor \times \left\lfloor \frac{n + 2p - f}{s} + 1 \right\rfloor$$

$$\frac{7 + 0 - 3}{2} + 1 = \frac{4}{2} + 1 = 3$$

## Summary of convolutions

$n \times n$  image       $f \times f$  filter

padding  $p$       stride  $s$

Output size:

$$\left\lfloor \frac{n+2p-f}{s} + 1 \right\rfloor \times \left\lfloor \underbrace{\frac{n+2p-f}{s}} + 1 \right\rfloor$$

# Technical note on cross-correlation vs. convolution


Convolution in math textbook:

2 <sup>7</sup>	3 <sup>2</sup>	7 <sup>5</sup>	4	6	2
6 <sup>9</sup>	6 <sup>0</sup>	9 <sup>4</sup>	8	7	4
3 <sup>-1</sup>	4 <sup>1</sup>	8 <sup>3</sup>	3	8	9
7	8	3	6	6	3
4	2	1	8	3	4
3	2	4	1	9	8

	3	4	5
	1	0	2
	-1	9	7

\*

7	2	5
9	0	4
-1	1	3



$$(A * B) * C = A * (B * C)$$



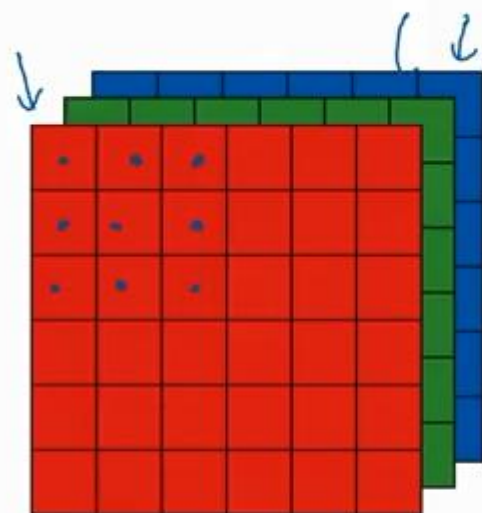
deeplearning.ai

# Convolutional Neural Networks

---

## Convolutions over volumes

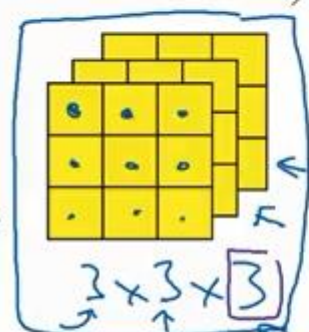
# Convolutions on RGB image



$6 \times 6 \times 3$   
 $\uparrow \quad \uparrow \quad \uparrow$

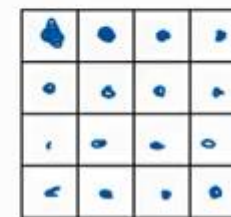


\*





$3 \times 3 \times 3$   
 27 numbers

=

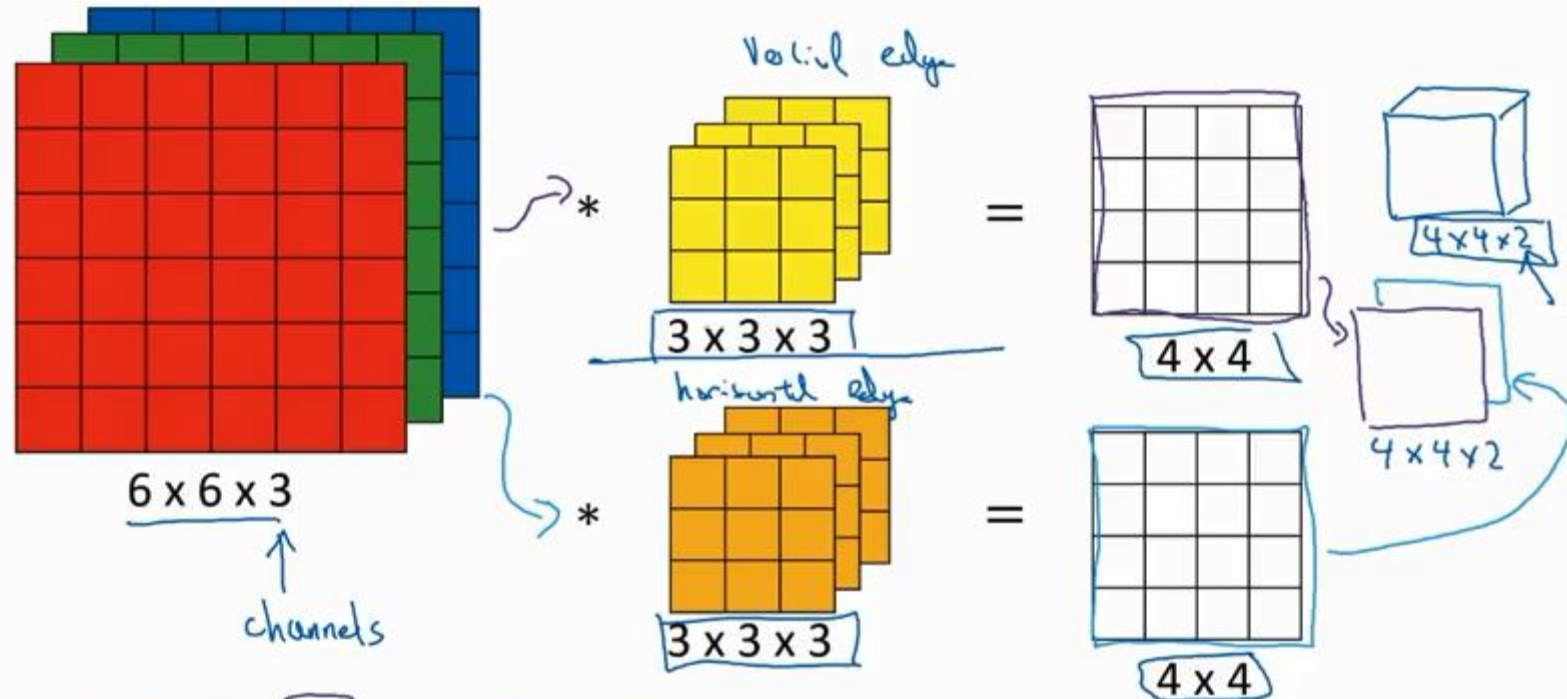


$4 \times 4$



Andrew Ng

# Multiple filters



Summary:  $n \times n \times n_c$   $*$   $f \times f \times n_c$   $\rightarrow$   $\frac{n-f+1}{4} \times \frac{n-f+1}{4} \times \frac{n_c'}{2}$

$6 \times 6 \times 3$   $3 \times 3 \times 3$   $4 \times 4 \times 2$   $\uparrow$  #filters



deeplearning.ai

# Convolutional Neural Networks


---

One layer of a  
convolutional  
network

---




# Example of a layer





## Number of parameters in one layer

If you have 10 filters that are  $3 \times 3 \times 3$  in one layer of a neural network, how many parameters does that layer have?



## Summary of notation

If layer  $l$  is a convolution layer:

$f^{[l]}$  = filter size

$p^{[l]}$  = padding

$s^{[l]}$  = stride

$n_c^{[l]}$  = number of filters

→ Each filter is:  $f^{[1]} \times f^{[2]} \times n_c^{[l]}$

Activations:  $a^{[l]} \rightarrow n_H^{[l]} \times n_W^{[l]} \times n_c^{[l]}$

Weights:  $f^{[1]} \times f^{[2]} \times n_c^{[l-1]} \times n_c^{[l]}$

bias:  $n_c^{[l]} = (1, 1, 1, n_c^{[l]})$  ← #filters in layer  $l$ .

Input:  $n_H^{[l-1]} \times n_W^{[l-1]} \times n_c^{[l-1]}$   
Output:  $n_H^{[l]} \times n_W^{[l]} \times n_c^{[l]}$

$$n_{HW}^{[l]} = \left\lfloor \frac{n_H^{[l-1]} + 2p^{[l]} - f^{[1]}}{s^{[1]}} + 1 \right\rfloor$$

$$A^{[l]} \rightarrow m \times \underbrace{n_H^{[l]} \times n_W^{[l]} \times n_c^{[l]}}_{n_c^{[l]} \times n_H^{[l]} \times n_W^{[l]}}$$




deeplearning.ai

# Convolutional Neural Networks


---

## A simple convolution network example

# Example ConvNet



## Types of layer in a convolutional network:

- Convolution (conv) ←
  - Pooling (pool) ←
  - Fully connected (FC) ←
- 




deeplearning.ai

# Convolutional Neural Networks

---

## Pooling layers


## Pooling layer: Max pooling



Hypoparameters:  
 $f = \underline{2}$   
 $s = \underline{2}$   
No parameters!



## Pooling layer: Max pooling





## Pooling layer: Average pooling

1	3	2	1
2	9	1	1
1	4	2	3
5	6	1	2

→

3.75	1.25
4	2

$$f=2$$

$$s=2$$

$$\underline{7 \times 7 \times 1000} \rightarrow 1 \times 1 \times 1000$$

# Summary of pooling

Hyperparameters:

f : filter size  
s : stride

$$\begin{aligned} f=2, s=2 \\ f=3, s=2 \end{aligned}$$

Max or average pooling

~~$\Rightarrow p$ : padding.~~

No parameters to learn!

$$\begin{aligned} & n_H \times n_W \times \underline{n_C} \\ & \quad \downarrow \\ & \left\lfloor \frac{n_H - f}{s} + 1 \right\rfloor \times \left\lfloor \frac{n_W - f}{s} + 1 \right\rfloor \\ & \quad \times \underline{n_C} \end{aligned}$$





deeplearning.ai

# Convolutional Neural Networks

---

## Convolutional neural network example

# Neural network example (LeNet-5)



CONV - POOL - CONV - POOL - FC - FC - FC - SOFTMAX

Andrew Ng

# Neural network example

	Activation shape	Activation Size	# parameters
Input:	(32,32,3)	— 3,072 $a^{[1]}$	0
CONV1 (f=5, s=1)	(28,28,8)	<u>6,272</u>	208 ←
POOL1	(14,14,8)	<u>1,568</u>	0 ←
CONV2 (f=5, s=1)	(10,10,16)	<u>1,600</u>	416 ←
POOL2	(5,5,16)	<u>400</u>	0 ←
FC3	(120,1)	<u>120</u>	48,001 }
FC4	(84,1)	<u>84</u>	10,081 }
Softmax	(10,1)	<u>10</u>	841





deeplearning.ai

# Convolutional Neural Networks

---

## Why convolutions?

# Why convolutions



$$5 \times 5 = 25$$


$$26$$

$$6 \times 26 = 156 \text{ parameters}$$

$$3,072 \times 4,704 \approx \underline{14M}$$



# Why convolutions




**Parameter sharing:** A feature detector (such as a vertical edge detector) that's useful in one part of the image is probably useful in another part of the image.

→ **Sparsity of connections:** In each layer, each output value depends only on a small number of inputs.



## Putting it together

Training set  $(x^{(1)}, y^{(1)}) \dots (x^{(m)}, y^{(m)})$ .



$$\text{Cost } J = \frac{1}{m} \sum_{i=1}^m \mathcal{L}(\hat{y}^{(i)}, y^{(i)})$$

Use gradient descent to optimize parameters to reduce  $J$