

Chapter 5: Object Detection



deeplearning.ai

Object Detection


Object
localization

What are localization and detection?

Image classification



Classification with
localization



Detection



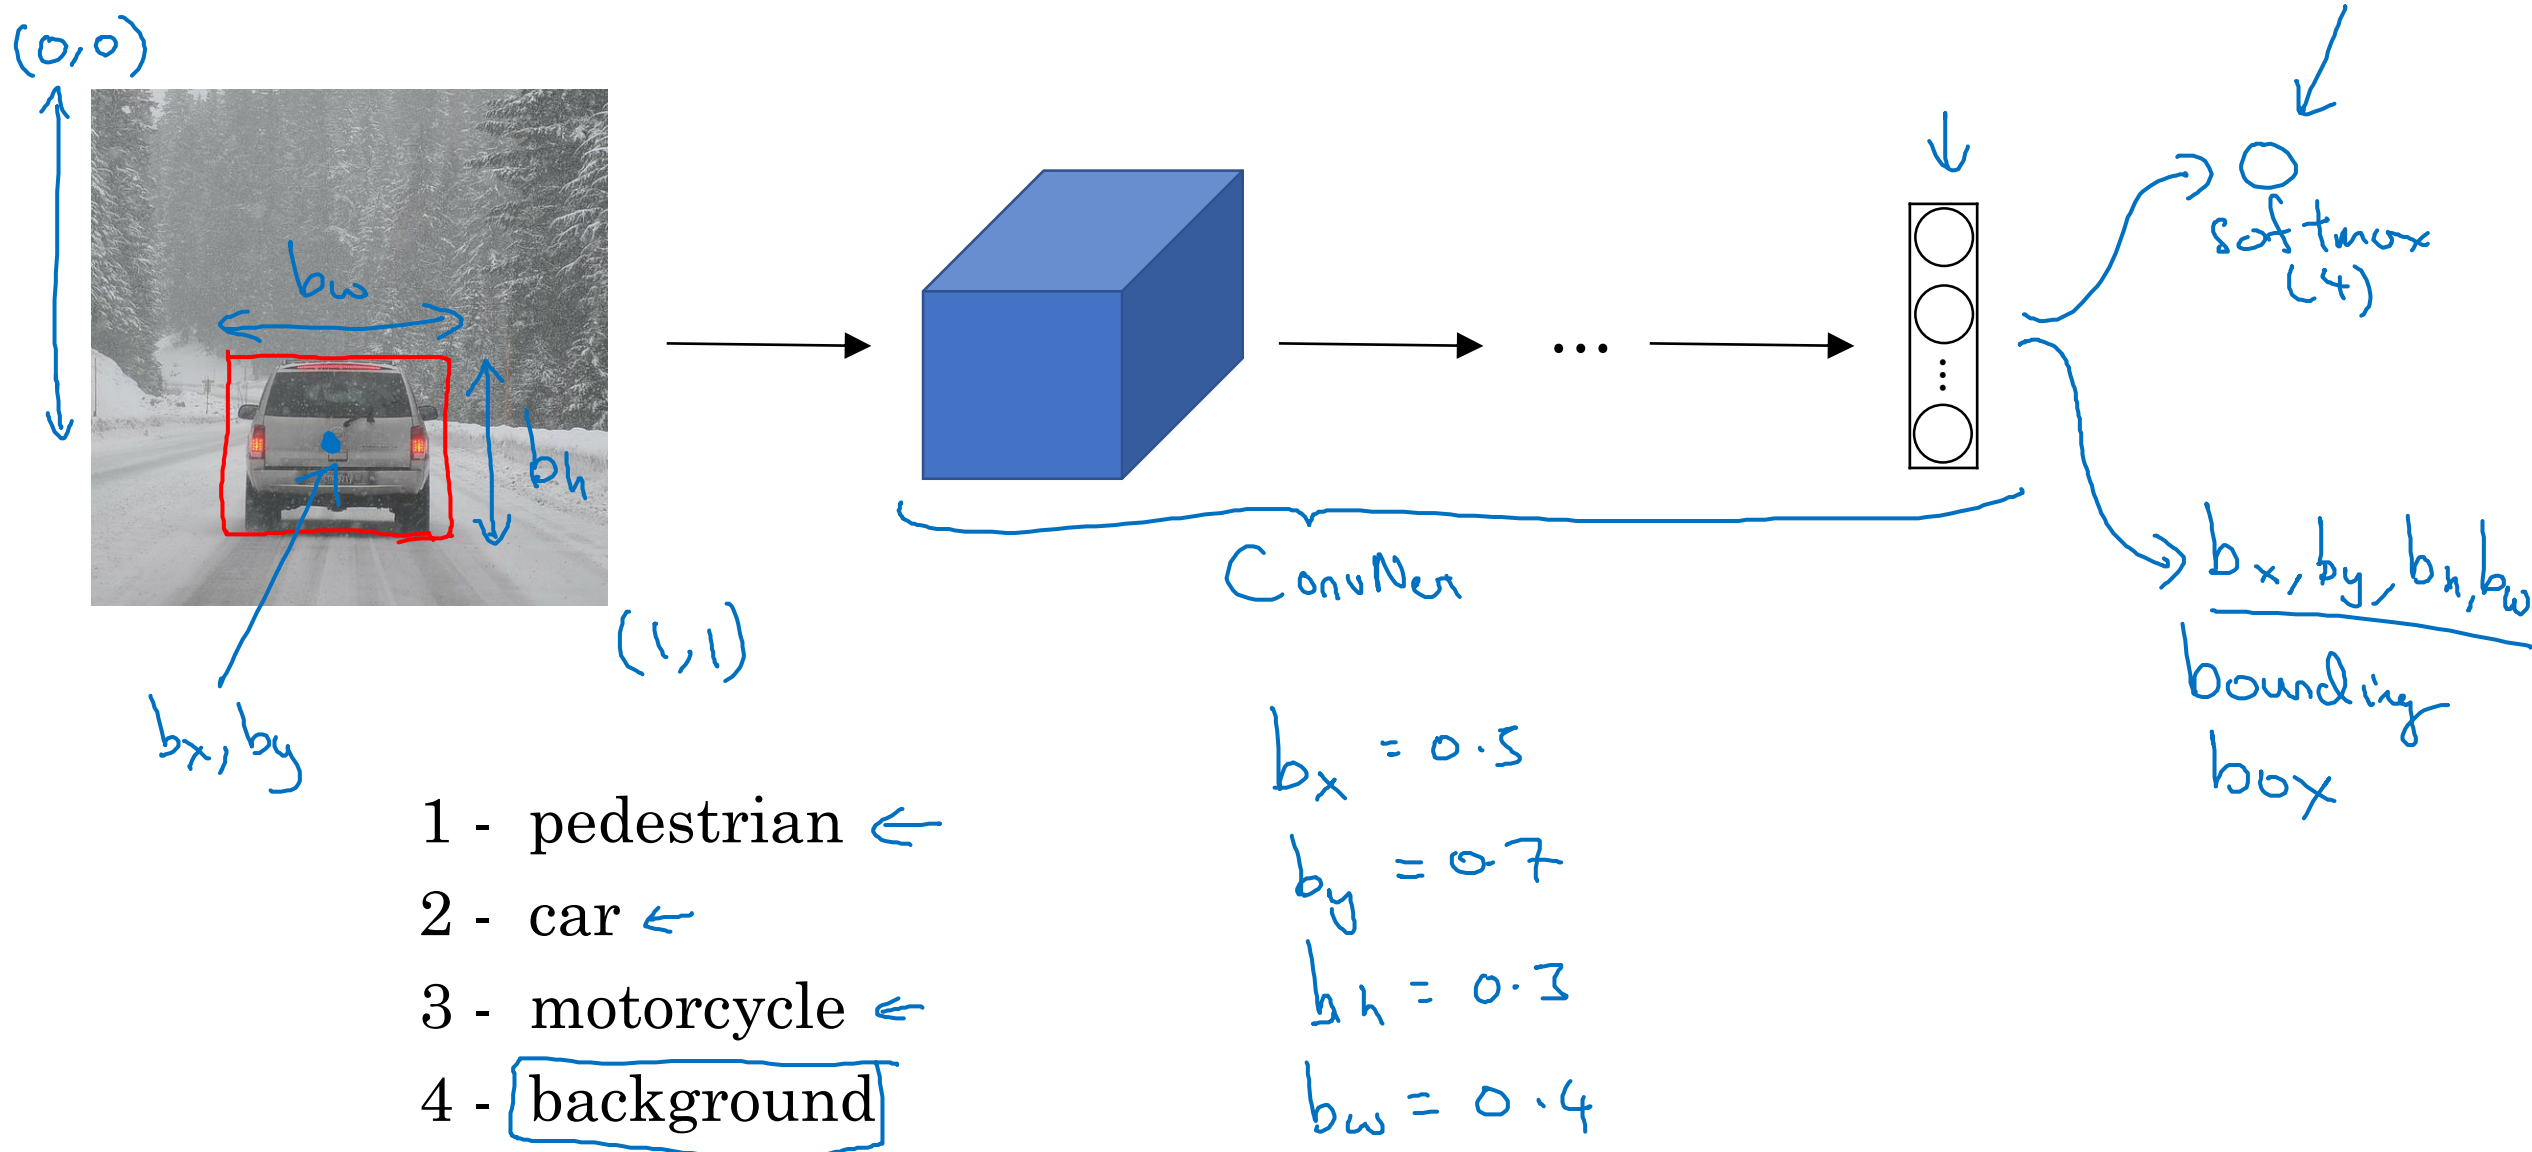
"Car"

"Car"

1 object

multiple
objects



Classification with localization



Defining the target label y

- 1 - pedestrian
- 2 - car ←
- 3 - motorcycle
- 4 - background ←

Need to output b_x, b_y, b_h, b_w , class label (1-4)



$$L(\hat{y}, y) = \begin{cases} (\hat{y}_1 - y_1)^2 + (\hat{y}_2 - y_2)^2 + \dots + (\hat{y}_8 - y_8)^2 & \text{if } \underline{y_1 = 1} \\ (\hat{y}_1 - y_1)^2 & \text{if } \underline{y_1 = 0} \end{cases}$$

$$\rightarrow y = \begin{bmatrix} p_c \\ b_x \\ b_y \\ b_h \\ b_w \\ c_1 \\ c_2 \\ c_3 \end{bmatrix} \quad \left. \begin{array}{l} \text{is there any} \\ \text{object?} \end{array} \right\}$$

(x, y)

$$\begin{bmatrix} 1 \\ b_x \\ b_y \\ b_h \\ b_w \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

$$\begin{bmatrix} 0 \\ \vdots \end{bmatrix} \quad \left. \begin{array}{l} p_c \\ \text{"don't care"} \end{array} \right\}$$





deeplearning.ai



Object Detection

Landmark
detection

Landmark detection



b_x, b_y, b_h, b_w



l1x, l1y, l2x, l2y, l3x, l3y, l4x, l4y, ..., l64x, l64y } X, y

l1x, l1y, ..., l32x, l32y



deeplearning.ai

Object Detection

Object
detection

Car detection example

Training set:

X

y



1



1



1



0





0





\rightarrow ConvNet $\rightarrow y$

Sliding windows detection


→ ConvNet → 0



→ ConvNet



Computation cost






deeplearning.ai


Object Detection

Convolutional
implementation of
sliding windows

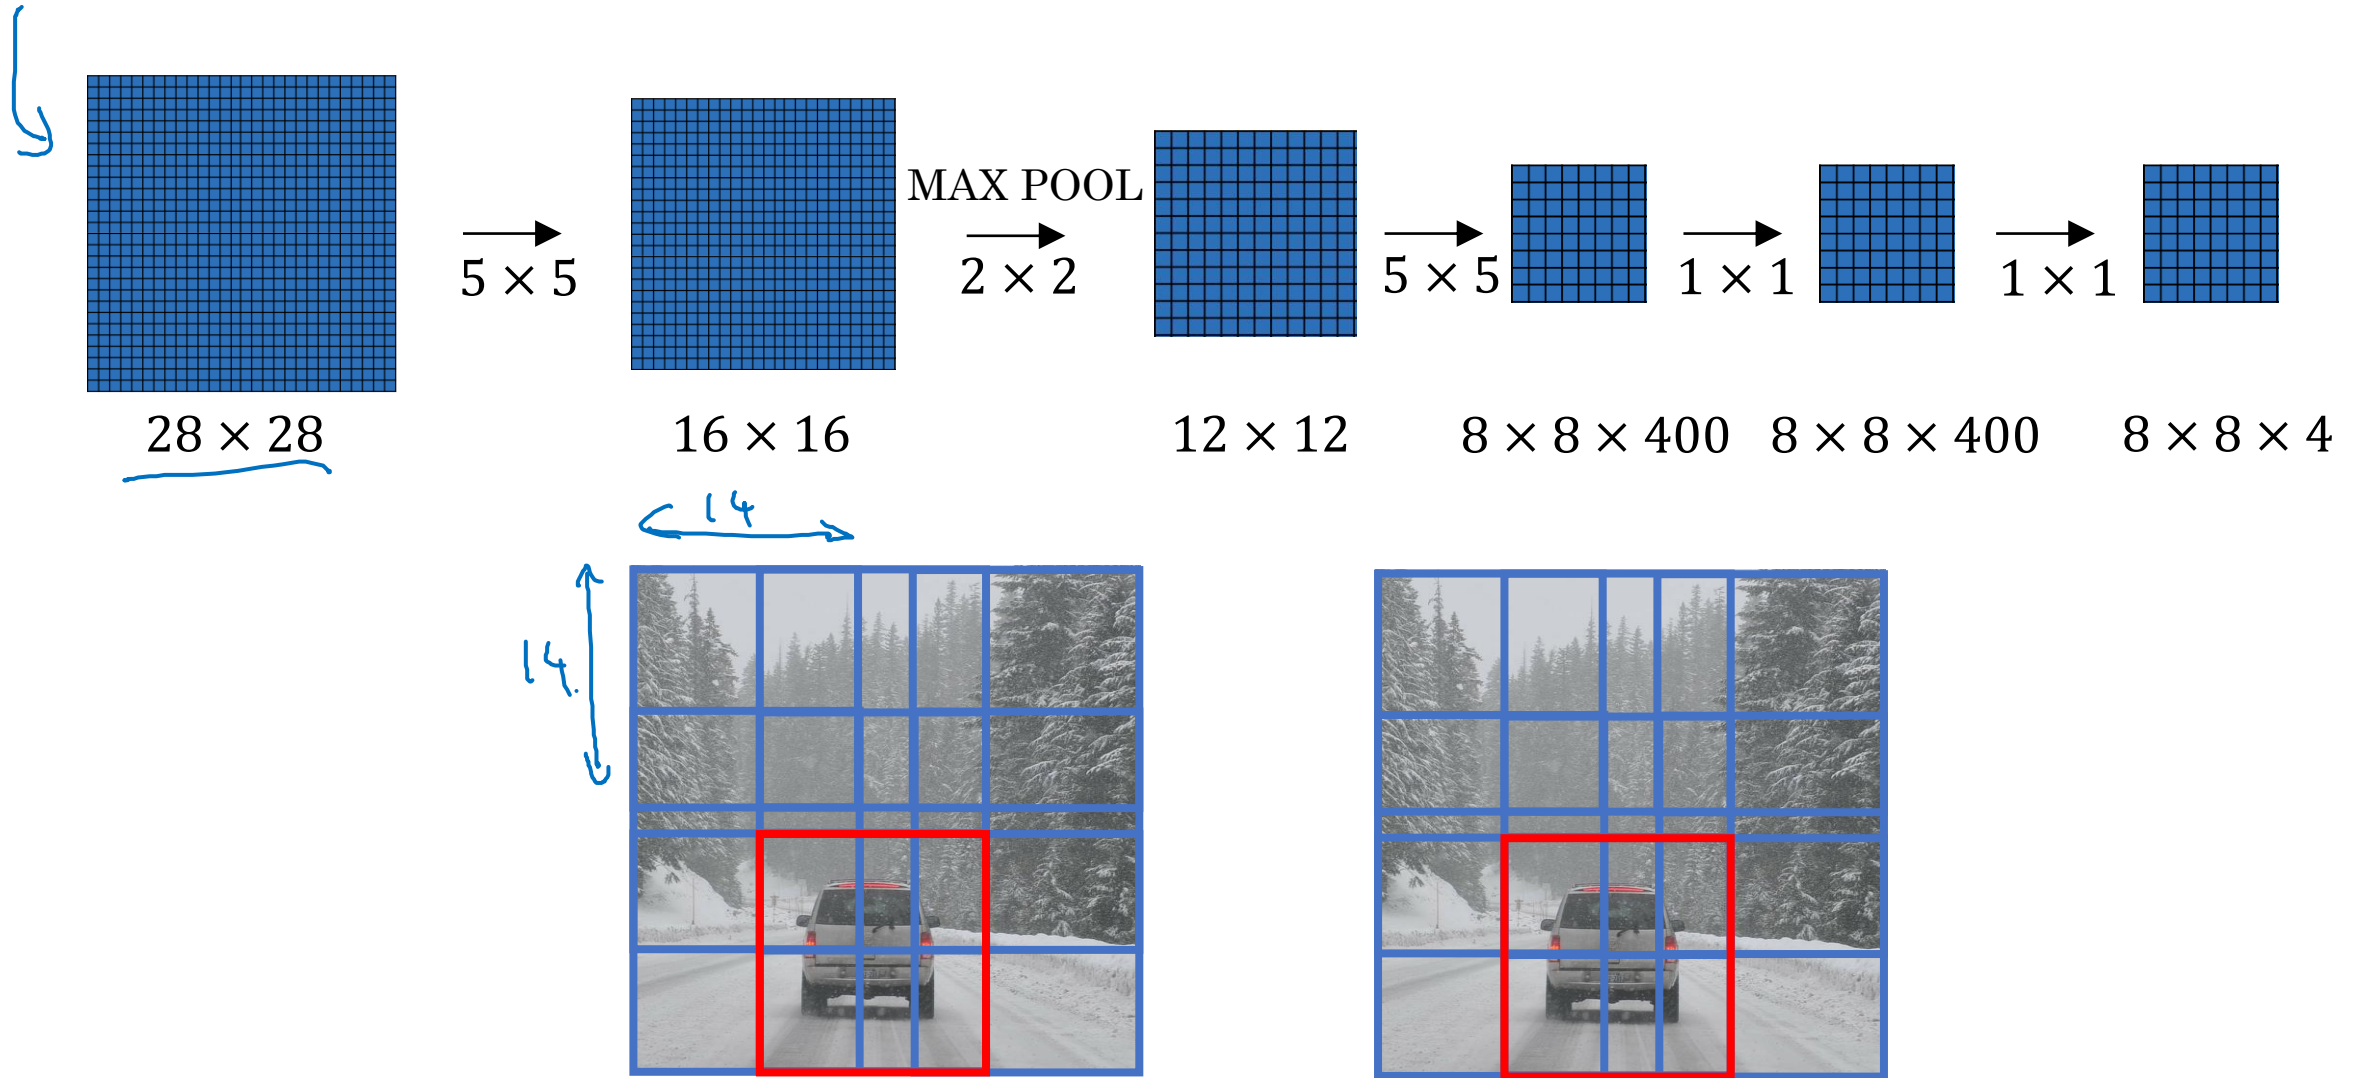
Turning FC layer into convolutional layers



Convolution implementation of sliding windows



Convolution implementation of sliding windows






deeplearning.ai

Object Detection

Intersection
over union

Evaluating object localization



Intersection over Union (IoU)

$$= \frac{\text{size of } \text{[yellow hatched box]}}{\text{size of } \text{[green hatched box]}}$$

“Correct” if $\text{IoU} \geq 0.5$ ←

0.6 ←

More generally, IoU is a measure of the overlap between two bounding boxes.



deeplearning.ai

Object Detection

Non-max
suppression

Non-max suppression example




Non-max suppression example




19x19

Non-max suppression example




P_c

Non-max suppression algorithm



19 × 19

Each output prediction is:



Discard all boxes with $p_c \leq 0.6$

→ While there are any remaining boxes:

- Pick the box with the largest p_c
Output that as a prediction.
- Discard any remaining box with $\text{IoU} \geq 0.5$ with the box output in the previous step




deeplearning.ai

Object Detection

Anchor boxes


Overlapping objects:




$$y = \begin{bmatrix} p_c \\ b_x \\ b_y \\ b_h \\ b_w \\ c_1 \\ c_2 \\ c_3 \end{bmatrix}$$

Annotations: A green arrow points from p_c to the center dot in the image. A blue arrow points from b_x, b_y to the top-left corner of the bounding box. A blue bracket groups c_1, c_2, c_3 .


Anchor box 1:



Anchor box 2:



$y =$




Anchor box algorithm

Previously:

Each object in training image is assigned to grid cell that contains that object's midpoint.

Output y :
 $3 \times 3 \times 8$




With two anchor boxes:

Each object in training image is assigned to grid cell that contains object's midpoint and anchor box for the grid cell with highest IoU.



(grid cell, anchor box)

Output y :
 $3 \times 3 \times 16$
 $3 \times 3 \times 2 \times 8$

Anchor box example



Anchor box 1: Anchor box 2:



y =

$$\begin{bmatrix} p_c \\ b_x \\ b_y \\ b_h \\ b_w \\ c_1 \\ c_2 \\ c_3 \\ p_c \\ b_x \\ b_y \\ b_h \\ b_w \\ c_1 \\ c_2 \\ c_3 \end{bmatrix}$$

1

bx

by

bh

bw

1

0

0

1

bx

by

bh

bw

0

1

0

car only?

0

?

?

?

?

?

?

?

1

bx

by

bh

bw

0


1

0

anchor box 1

anchor box 2

YOLO — You Only Look Once



YOLO — You Only Look Once

- Original paper (CVPR 2016. OpenCV People's Choice Award) <https://arxiv.org/pdf/1506.02640v5.pdf>
- YOLOv2: <https://arxiv.org/pdf/1612.08242v1.pdf>
- YOLOv3: <https://arxiv.org/pdf/1804.02767.pdf>
- Open source: <https://pjreddie.com/darknet/yolo/>